

```
In [2]: import matplotlib.pyplot as plt
from PIL import Image
import seaborn as sns
import numpy as np
import pandas as pd
import os
from tensorflow.keras.utils import to_categorical
from glob import glob

from sklearn.model_selection import train_test_split
import keras
from keras.models import Sequential
from keras.layers import Dense, Dropout
import tensorflow as tf
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import ReduceLROnPlateau
from tensorflow.keras import layers
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten

from sklearn.metrics import confusion_matrix
import itertools
import matplotlib.pyplot as plt

import warnings
warnings.filterwarnings("ignore")
```

In [2]: `pip install seaborn`

```
Requirement already satisfied: seaborn in c:\programdata\anaconda3\lib\site-packages (0.11.2)
Requirement already satisfied: matplotlib>=2.2 in c:\programdata\anaconda3\lib\site-packages (from seaborn) (3.5.1)
Requirement already satisfied: numpy>=1.15 in c:\programdata\anaconda3\lib\site-packages (from seaborn) (1.21.5)
Requirement already satisfied: scipy>=1.0 in c:\programdata\anaconda3\lib\site-packages (from seaborn) (1.7.3)
Requirement already satisfied: pandas>=0.23 in c:\programdata\anaconda3\lib\site-packages (from seaborn) (1.4.2)
Requirement already satisfied: pillow>=6.2.0 in c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.2->seaborn) (9.0.1)
Requirement already satisfied: packaging>=20.0 in c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.2->seaborn) (21.3)
Requirement already satisfied: pyparsing>=2.2.1 in c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.2->seaborn) (3.0.4)
Requirement already satisfied: python-dateutil>=2.7 in c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.2->seaborn) (2.8.2)
Requirement already satisfied: fonttools>=4.22.0 in c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.2->seaborn) (4.25.0)
Requirement already satisfied: cycler>=0.10 in c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.2->seaborn) (0.11.0)
Requirement already satisfied: kiwisolver>=1.0.1 in c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.2->seaborn) (1.3.2)
Requirement already satisfied: pytz>=2020.1 in c:\programdata\anaconda3\lib\site-packages (from pandas>=0.23->seaborn) (2021.3)
Requirement already satisfied: six>=1.5 in c:\programdata\anaconda3\lib\site-packages (from python-dateutil>=2.7->matplotlib>=2.2->seaborn) (1.16.0)
Note: you may need to restart the kernel to use updated packages.
```

In [3]: `pip install pydotplus`

```
Requirement already satisfied: pydotplus in c:\programdata\anaconda3\lib\site-packages (2.0.2)
Note: you may need to restart the kernel to use updated packages.
Requirement already satisfied: pyparsing>=2.0.1 in c:\programdata\anaconda3\lib\site-packages (from pydotplus) (3.0.4)
```

```
In [3]: df = pd.read_csv('skinCancer/HAM10000_metadata.csv')
df.head(10)
```

Out[3]:

	lesion_id	image_id	dx	dx_type	age	sex	localization
0	HAM_0000118	ISIC_0027419	bkl	histo	80.0	male	scalp
1	HAM_0000118	ISIC_0025030	bkl	histo	80.0	male	scalp
2	HAM_0002730	ISIC_0026769	bkl	histo	80.0	male	scalp
3	HAM_0002730	ISIC_0025661	bkl	histo	80.0	male	scalp
4	HAM_0001466	ISIC_0031633	bkl	histo	75.0	male	ear
5	HAM_0001466	ISIC_0027850	bkl	histo	75.0	male	ear
6	HAM_0002761	ISIC_0029176	bkl	histo	60.0	male	face
7	HAM_0002761	ISIC_0029068	bkl	histo	60.0	male	face
8	HAM_0005132	ISIC_0025837	bkl	histo	70.0	female	back
9	HAM_0005132	ISIC_0025209	bkl	histo	70.0	female	back

```
In [4]: df.isnull().sum()
```

```
Out[4]: lesion_id      0
image_id      0
dx            0
dx_type      0
age          57
sex          0
localization  0
dtype: int64
```

```
In [5]: df.count()
```

```
Out[5]: lesion_id      10015
image_id      10015
dx            10015
dx_type      10015
age          9958
sex          10015
localization  10015
dtype: int64
```

```
In [6]: df['age'].fillna(int(df['age'].mean()),inplace=True)
df.isnull().sum()
```

```
Out[6]: lesion_id      0
image_id      0
dx            0
dx_type       0
age           0
sex           0
localization  0
dtype: int64
```

```
In [7]: lesion_type_dict = {
        'nv': 'Melanocytic nevi',
        'mel': 'Melanoma',
        'bkl': 'Benign keratosis-like lesions ',
        'bcc': 'Basal cell carcinoma',
        'akiec': 'Actinic keratoses',
        'vasc': 'Vascular lesions',
        'df': 'Dermatofibroma'
    }
base_skin_dir = 'skinCancer'

imageid_path_dict = {os.path.splitext(os.path.basename(x))[0]: x
                      for x in glob(os.path.join(base_skin_dir, '*', '*.jpg'))}
```

```
In [8]: df['path'] = df['image_id'].map(imageid_path_dict.get)
df['cell_type'] = df['dx'].map(lesion_type_dict.get)
df['cell_type_idx'] = pd.Categorical(df['cell_type']).codes
df.head()
```

```
Out[8]:
```

	lesion_id	image_id	dx	dx_type	age	sex	localization	
0	HAM_0000118	ISIC_0027419	bkl	histo	80.0	male	scalp	skinCancer\HAM10000_images_
1	HAM_0000118	ISIC_0025030	bkl	histo	80.0	male	scalp	skinCancer\HAM10000_images_
2	HAM_0002730	ISIC_0026769	bkl	histo	80.0	male	scalp	skinCancer\HAM10000_images_
3	HAM_0002730	ISIC_0025661	bkl	histo	80.0	male	scalp	skinCancer\HAM10000_images_
4	HAM_0001466	ISIC_0031633	bkl	histo	75.0	male	ear	skinCancer\HAM10000_images_

```
In [9]: df['image'] = df['path'].map(lambda x: np.asarray(Image.open(x).resize((125,100)))
```

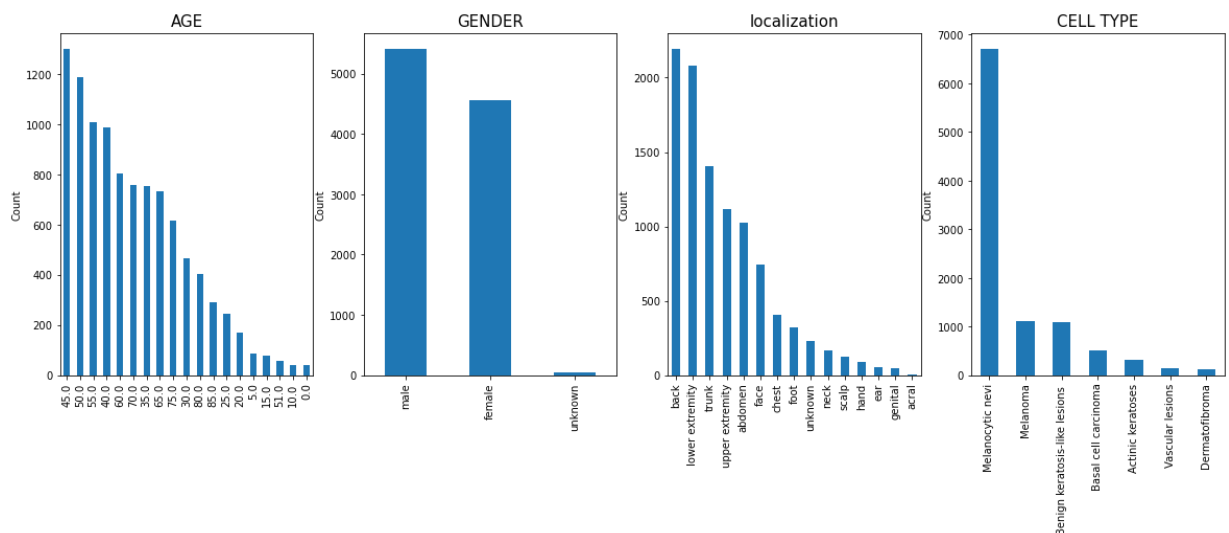
```
In [10]: plt.figure(figsize=(20,10))
plt.subplots_adjust(left=0.125, bottom=1, right=0.9, top=2, hspace=0.2)
plt.subplot(2,4,1)
plt.title("AGE", fontsize=15)
plt.ylabel("Count")
df['age'].value_counts().plot.bar()

plt.subplot(2,4,2)
plt.title("GENDER", fontsize=15)
plt.ylabel("Count")
df['sex'].value_counts().plot.bar()


plt.subplot(2,4,3)
plt.title("localization", fontsize=15)
plt.ylabel("Count")
plt.xticks(rotation=45)
df['localization'].value_counts().plot.bar()

plt.subplot(2,4,4)
plt.title("CELL TYPE", fontsize=15)
plt.ylabel("Count")
df['cell_type'].value_counts().plot.bar()
```

```
Out[10]: <AxesSubplot:title={'center':'CELL TYPE'}, ylabel='Count'>
```



```
In [11]: plt.figure(figsize=(15,10))
plt.subplot(1,2,1)
df['dx'].value_counts().plot.pie(autopct="%1.1f%%")
plt.subplot(1,2,2)
df['dx_type'].value_counts().plot.pie(autopct="%1.1f%%")
plt.show()
```



```
In [12]: features=df.drop(columns=['cell_type_idx','image_id'],axis=1)
target=df['cell_type_idx']
features.head()
```

Out[12]:

	lesion_id	dx	dx_type	age	sex	localization	
0	HAM_0000118	bkl	histo	80.0	male	scalp	skinCancer\HAM10000_images_part_1\ISIC_002
1	HAM_0000118	bkl	histo	80.0	male	scalp	skinCancer\HAM10000_images_part_1\ISIC_002
2	HAM_0002730	bkl	histo	80.0	male	scalp	skinCancer\HAM10000_images_part_1\ISIC_002
3	HAM_0002730	bkl	histo	80.0	male	scalp	skinCancer\HAM10000_images_part_1\ISIC_002
4	HAM_0001466	bkl	histo	75.0	male	ear	skinCancer\HAM10000_images_part_2\ISIC_003

```
In [13]: x_train_o, x_test_o, y_train_o, y_test_o = train_test_split(features, target, test_size=0.2,
tf.unique(x_train_o.cell_type.values)
```

```
Out[13]: Unique(y=<tf.Tensor: shape=(7,), dtype=string, numpy=
array([b'Melanocytic nevi', b'Basal cell carcinoma', b'Melanoma',
      b'Vascular lesions', b'Benign keratosis-like lesions ',
      b'Actinic keratoses', b'Dermatofibroma'], dtype=object)>, idx=<tf.Tensor:
shape=(7511,), dtype=int32, numpy=array([0, 1, 0, ..., 1, 0, 0])>)
```

```
In [14]: x_train = np.asarray(x_train_o['image'].tolist())
x_test = np.asarray(x_test_o['image'].tolist())

x_train_mean = np.mean(x_train)
x_train_std = np.std(x_train)

x_test_mean = np.mean(x_test)
x_test_std = np.std(x_test)

x_train = (x_train - x_train_mean)/x_train_std
x_test = (x_test - x_test_mean)/x_test_std
```

```
In [15]: # Perform one-hot encoding on the labels
y_train = to_categorical(y_train_o, num_classes = 7)
y_test = to_categorical(y_test_o, num_classes = 7)
y_test
```

```
Out[15]: array([[0., 0., 0., ..., 0., 1., 0.],
               [0., 0., 0., ..., 1., 0., 0.],
               [1., 0., 0., ..., 0., 0., 0.],
               ...,
               [0., 0., 0., ..., 1., 0., 0.],
               [0., 0., 0., ..., 1., 0., 0.],
               [0., 0., 0., ..., 1., 0., 0.]], dtype=float32)
```

```
In [17]: y_test[1]
```

```
Out[17]: array([0., 0., 0., 0., 1., 0., 0.], dtype=float32)
```

```
In [18]: y_train
```

```
Out[18]: array([[0., 0., 0., ..., 1., 0., 0.],
               [0., 1., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 1., 0., 0.],
               ...,
               [0., 1., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 1., 0., 0.],
               [0., 0., 0., ..., 1., 0., 0.]], dtype=float32)
```


In [19]: `y_test_o.value_counts()`

Out[19]:

4	1693
2	277
5	274
1	123
0	68
6	39
3	30

Name: cell_type_idx, dtype: int64

In [16]: `# x_train, x_validate, y_train, y_validate = train_test_split(x_train, y_train, t`

`# Reshape image in 3 dimensions (height = 100, width = 125 , canal = 3)`

`x_train = x_train.reshape(x_train_o.shape[0], *(100, 125, 3))`

`x_test = x_test.reshape(x_test_o.shape[0], *(100, 125, 3))`

In [17]: `x_train = x_train.reshape(x_train.shape[0],125*100*3)`

`x_test = x_test.reshape(x_test.shape[0],125*100*3)`

`print(x_train.shape)`

`print(x_test.shape)`

(7511, 37500)

(2504, 37500)

In [22]: `print(x_train)`

```
[ [ 1.55231607e+00 -3.87567404e-01 -1.72024796e-01 ... 9.27242505e-01
   -5.60001490e-01 -6.03110012e-01]
  [ 3.23723203e-01 -4.09121664e-01 -1.72024796e-01 ... 3.66831724e-01
   -3.44458882e-01 -1.07362013e-01]
  [ 1.22900216e+00 -7.97098359e-01 -2.15133317e-01 ... 1.09967659e+00
   -3.87567404e-01 -2.15133317e-01]
  ...
  [-2.56454774e+00 -2.90941592e+00 -2.43522218e+00 ... -2.78009035e+00
   -2.99563296e+00 -2.65076479e+00]
  [ 4.53048767e-01 -1.07362013e-01 -8.58077525e-02 ... 3.88385985e-01
   -4.26992309e-02 4.09290704e-04]
  [ 1.63853311e+00 -3.01350360e-01 -2.15133317e-01 ... 1.09967659e+00
   -4.52230186e-01 -6.46218533e-01]]
```

Decision Tree

```

In [23]: depth = range(5,51,5)
testing_accuracy = []
training_accuracy = []
score = 0

for i in depth:
    tree = DecisionTreeClassifier(max_depth = i, criterion = 'entropy')
    tree.fit(x_train, y_train)

    y_predict_train = tree.predict(x_train)
    training_accuracy.append(accuracy_score(y_train, y_predict_train))

    y_predict_test = tree.predict(x_test)
    acc_score = accuracy_score(y_test, y_predict_test)
    testing_accuracy.append(acc_score)

    print(i)


    if score < acc_score:
        score = acc_score
        best_depth = i

sns.lineplot(depth, training_accuracy)
sns.scatterplot(depth, training_accuracy)
sns.lineplot(depth, testing_accuracy)
sns.scatterplot(depth, testing_accuracy)
plt.legend(['training accuracy', 'testing accuracy'])

```

5
10
15
20
25
30
35
40
45
50

Out[23]: <matplotlib.legend.Legend at 0x18880032ee0>



```
In [24]: print('This is the best depth for Decision Tree Classifier: ', best_depth, '\nAcc
result = confusion_matrix(y_test.argmax(axis=1), y_predict_test.argmax(axis=1),)
print("Confusion Matrix:", "\n", result)
report = classification_report(y_test.argmax(axis=1), y_predict_test.argmax(axis=
print("Classification Report:" , "\n", report)
```

This is the best depth for Decision Tree Classifier: 30

Accuracy score is: 0.6285942492012779

Confusion Matrix:

```
[[ 17  12  13   1  20   5   0]
 [ 17  26  30   5  33   6   6]
 [ 12  19  90   9 100  45   2]
 [  3   4   6   1  10   5   1]
 [ 12  35 116  10 135 150  20]
 [ 12  16  37   2 118  83   6]
 [  2   6   5   1  19   4   2]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.23	0.25	0.24	68
1	0.22	0.21	0.22	123
2	0.30	0.32	0.31	277
3	0.03	0.03	0.03	30
4	0.82	0.80	0.81	1693
5	0.28	0.30	0.29	274
6	0.05	0.05	0.05	39
accuracy			0.63	2504
macro avg	0.28	0.28	0.28	2504
weighted avg	0.64	0.63	0.63	2504

Random Forest

```
In [ ]: #random Forest
depth = range(5,51,5)
testing_accuracy = []
training_accuracy = []
score = 0

for i in depth:
    tree = RandomForestClassifier(max_depth = i, criterion = 'gini', random_state=42)
    tree.fit(x_train, y_train)

    y_predict_train = tree.predict(x_train)
    training_accuracy.append(accuracy_score(y_train, y_predict_train))

    y_predict_test = tree.predict(x_test)
    acc_score = accuracy_score(y_test,y_predict_test)
    testing_accuracy.append(acc_score)

    print(i)

    if score < acc_score:
        score = acc_score
        best_depth = i

sns.lineplot(depth, training_accuracy)
sns.scatterplot(depth, training_accuracy)
sns.lineplot(depth, testing_accuracy)
sns.scatterplot(depth, testing_accuracy)
plt.legend(['training accuracy', 'testing accuracy'])
```

```
In [26]: print('This is the best depth for Decision Tree Classifier: ', best_depth, '\nAcc
result = confusion_matrix(y_test.argmax(axis=1), y_predict_test.argmax(axis=1),)
print("Confusion Matrix:", "\n", result)
report = classification_report(y_test.argmax(axis=1), y_predict_test.argmax(axis=
print("Classification Report:" , "\n", report)
```

This is the best depth for Decision Tree Classifier: 10

Accuracy score is: 0.6202076677316294

Confusion Matrix:

```
[[ 63   1   0   0   4   0   0]
 [ 99   7   2   0  15   0   0]
 [164   0  25   0  88   0   0]
 [ 23   0   0   0   7   0   0]
 [221   0  10   0 1459   3   0]
 [149   0   3   0  114   8   0]
 [ 21   0   0   0  18   0   0]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.09	0.93	0.16	68
1	0.88	0.06	0.11	123
2	0.62	0.09	0.16	277
3	0.00	0.00	0.00	30
4	0.86	0.86	0.86	1693
5	0.73	0.03	0.06	274
6	0.00	0.00	0.00	39
accuracy			0.62	2504
macro avg	0.45	0.28	0.19	2504
weighted avg	0.77	0.62	0.61	2504

KNN

```
In [27]: k = range(1,102,3)
testing_accuracy = []
training_accuracy = []
score = 0

for i in k:
    knn = KNeighborsClassifier(n_neighbors = i)
    knn.fit(x_train, y_train)

    y_predict_train = knn.predict(x_train)
    training_accuracy.append(accuracy_score(y_train, y_predict_train))

    y_predict_test = knn.predict(x_test)
    acc_score = accuracy_score(y_test,y_predict_test)
    testing_accuracy.append(acc_score)

    print(i)


    if score < acc_score:
        score = acc_score
        best_k = i

sns.lineplot(k, training_accuracy)
sns.scatterplot(k, training_accuracy)
sns.lineplot(k, testing_accuracy)
sns.scatterplot(k, testing_accuracy)
plt.legend(['training accuracy', 'testing accuracy'])
```

1
4
7
10
13
16
19
22
25
28
31
34
37
40
43
46
49
52
55
58
61
64
67
70
73
76
79

82
85
88
91
94
97
100

Out[27]: <matplotlib.legend.Legend at 0x188f40ca340>



```
In [28]: print('This is the best depth for Decision Tree Classifier: ', best_depth, '\nAcc
result = confusion_matrix(y_test.argmax(axis=1), y_predict_test.argmax(axis=1),)
print("Confusion Matrix:", "\n", result)
report = classification_report(y_test.argmax(axis=1), y_predict_test.argmax(axis=
print("Classification Report:" , "\n", report)
```

This is the best depth for Decision Tree Classifier: 10

Accuracy score is: 0.6793130990415336

Confusion Matrix:

```
[[ 50   0   0   0  18   0   0]
 [ 77   0   1   0  45   0   0]
 [131   0   6   0 140   0   0]
 [ 13   0   1   0  16   0   0]
 [121   0   2   0 1570  0   0]
 [ 82   0   5   0 187   0   0]
 [ 13   0   0   0  26   0   0]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.10	0.74	0.18	68
1	0.00	0.00	0.00	123
2	0.40	0.02	0.04	277
3	0.00	0.00	0.00	30
4	0.78	0.93	0.85	1693
5	0.00	0.00	0.00	274
6	0.00	0.00	0.00	39
accuracy			0.65	2504
macro avg	0.18	0.24	0.15	2504
weighted avg	0.58	0.65	0.58	2504

```
In [29]: x_train.shape
```

```
Out[29]: (7511, 37500)
```

MLP


```
In [30]: # define the keras model
model = Sequential()

model.add(Dense(units= 64, kernel_initializer = 'uniform', activation = 'relu', i
model.add(Dense(units= 128, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dense(units= 64, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dense(units= 32, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dense(units = 7, kernel_initializer = 'uniform', activation = 'softmax'
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 64)	240064
dense_1 (Dense)	(None, 128)	8320
dense_2 (Dense)	(None, 64)	8256
dense_3 (Dense)	(None, 32)	2080
dense_4 (Dense)	(None, 7)	231
=====		
Total params: 2,418,951		
Trainable params: 2,418,951		
Non-trainable params: 0		

```
In [31]: # compile the keras model
model.compile(optimizer = "Adam", loss = 'categorical_crossentropy', metrics = [

# fit the keras model on the dataset
history = model.fit(x_train, y_train, batch_size = 20, epochs = 25)
```

```
Epoch 1/25
376/376 [=====] - 9s 16ms/step - loss: 0.9923 - accuracy: 0.6678
Epoch 2/25
376/376 [=====] - 6s 16ms/step - loss: 0.8973 - accuracy: 0.6865
Epoch 3/25
376/376 [=====] - 6s 16ms/step - loss: 0.8596 - accuracy: 0.7000
Epoch 4/25
376/376 [=====] - 6s 16ms/step - loss: 0.8178 - accuracy: 0.7092
Epoch 5/25
376/376 [=====] - 6s 17ms/step - loss: 0.7929 - accuracy: 0.7131
Epoch 6/25
376/376 [=====] - 6s 17ms/step - loss: 0.7731 - accuracy: 0.7256
Epoch 7/25
376/376 [=====] - 6s 17ms/step - loss: 0.7331 - accuracy: 0.7351
Epoch 8/25
376/376 [=====] - 6s 17ms/step - loss: 0.7064 - accuracy: 0.7383
Epoch 9/25
376/376 [=====] - 6s 17ms/step - loss: 0.6761 - accuracy: 0.7556
Epoch 10/25
376/376 [=====] - 6s 16ms/step - loss: 0.6542 - accuracy: 0.7585
Epoch 11/25
376/376 [=====] - 6s 16ms/step - loss: 0.6337 - accuracy: 0.7662
Epoch 12/25
376/376 [=====] - 6s 17ms/step - loss: 0.6045 - accuracy: 0.7759
Epoch 13/25
376/376 [=====] - 6s 17ms/step - loss: 0.5875 - accuracy: 0.7835
Epoch 14/25
376/376 [=====] - 6s 17ms/step - loss: 0.5539 - accuracy: 0.7988
Epoch 15/25
376/376 [=====] - 6s 17ms/step - loss: 0.5272 - accuracy: 0.8083
Epoch 16/25
376/376 [=====] - 6s 17ms/step - loss: 0.5191 - accuracy: 0.8083
Epoch 17/25
376/376 [=====] - 6s 16ms/step - loss: 0.5057 - accuracy: 0.8129
```

```

Epoch 18/25
376/376 [=====] - 6s 16ms/step - loss: 0.4894 - accuracy: 0.8223
Epoch 19/25
376/376 [=====] - 6s 16ms/step - loss: 0.4619 - accuracy: 0.8304
Epoch 20/25
376/376 [=====] - 6s 16ms/step - loss: 0.4364 - accuracy: 0.8386
Epoch 21/25
376/376 [=====] - 6s 16ms/step - loss: 0.4161 - accuracy: 0.8484
Epoch 22/25
376/376 [=====] - 6s 16ms/step - loss: 0.4188 - accuracy: 0.8466
Epoch 23/25
376/376 [=====] - 6s 16ms/step - loss: 0.3984 - accuracy: 0.8518
Epoch 24/25
376/376 [=====] - 6s 16ms/step - loss: 0.3836 - accuracy: 0.8662
Epoch 25/25
376/376 [=====] - 6s 16ms/step - loss: 0.3669 - accuracy: 0.8655

```

```

In [32]: accuracy = model.evaluate(x_test, y_test, verbose=1)[1]
print("Test: accuracy = ",accuracy*100,"%")

```

```

79/79 [=====] - 1s 9ms/step - loss: 1.1677 - accuracy: 0.7017
Test: accuracy = 70.16773223876953 %

```

CNN

```

In [18]: x_train = x_train.reshape(x_train.shape[0], 125,100,3)
x_test = x_test.reshape(x_test.shape[0], 125, 100, 3)
print(x_train.shape)
print(x_test.shape)

```

```

(7511, 125, 100, 3)
(2504, 125, 100, 3)

```

```
In [19]: input_shape = (125, 100, 3)
num_classes = 7

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', padding = 'Same', input_shape=input_shape))
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', padding = 'Same'))
model.add(MaxPooling2D(pool_size = (2, 2)))
model.add(Dropout(0.16))

model.add(Conv2D(64, (3, 3), activation='relu', padding = 'same'))
model.add(Conv2D(64, (3, 3), activation='relu', padding = 'Same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.20))

model.add(Flatten())
model.add(Dense(256, activation='relu'))
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.3))
model.add(Dense(num_classes, activation='softmax'))
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 125, 100, 32)	896
conv2d_1 (Conv2D)	(None, 125, 100, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 62, 50, 32)	0
dropout (Dropout)	(None, 62, 50, 32)	0
conv2d_2 (Conv2D)	(None, 62, 50, 64)	18496
conv2d_3 (Conv2D)	(None, 62, 50, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 31, 25, 64)	0
dropout_1 (Dropout)	(None, 31, 25, 64)	0
flatten (Flatten)	(None, 49600)	0
dense (Dense)	(None, 256)	12697856
dense_1 (Dense)	(None, 128)	32896
dropout_2 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 7)	903
=====		
Total params: 12,797,223		
Trainable params: 12,797,223		

Non-trainable params: 0

```
In [20]: model.compile(optimizer= "Adam",  
                        loss='categorical_crossentropy',  
                        metrics=['accuracy'])  
  
learning_rate_reduction = ReduceLROnPlateau(monitor='val_accuracy',  
                                              patience=4,  
                                              verbose=1,  
                                              factor=0.5,  
                                              min_lr=0.00001)
```

In [21]: `#x_train, x_validate, y_train, y_validate = train_test_split(x_train, y_train, te`

```
# Reshape image in 3 dimensions (height = 100, width = 125 , canal = 3)
x_train = x_train.reshape(x_train.shape[0], *(125, 100, 3))
x_test = x_test.reshape(x_test.shape[0], *(125,100, 3))
#x_validate = x_validate.reshape(x_validate.shape[0], *(100, 125, 3))
```

```
tf.config.run_functions_eagerly(True)
```

```
model.fit(x_train, y_train, batch_size = 100, epochs = 20)
```

Epoch 1/20

76/76 [=====] - 498s 7s/step - loss: 1.1379 - accuracy: 0.6457

Epoch 2/20

76/76 [=====] - 491s 6s/step - loss: 0.8668 - accuracy: 0.6867

Epoch 3/20

76/76 [=====] - 491s 6s/step - loss: 0.8072 - accuracy: 0.7047

Epoch 4/20

76/76 [=====] - 492s 6s/step - loss: 0.7610 - accuracy: 0.7251

Epoch 5/20

76/76 [=====] - 496s 7s/step - loss: 0.7106 - accuracy: 0.7386

Epoch 6/20

76/76 [=====] - 492s 6s/step - loss: 0.6676 - accuracy: 0.7565

Epoch 7/20

76/76 [=====] - 492s 6s/step - loss: 0.6182 - accuracy: 0.7747

Epoch 8/20

76/76 [=====] - 493s 6s/step - loss: 0.5940 - accuracy: 0.7842

Epoch 9/20

76/76 [=====] - 499s 7s/step - loss: 0.5232 - accuracy: 0.8103

Epoch 10/20

76/76 [=====] - 491s 6s/step - loss: 0.4530 - accuracy: 0.8333

Epoch 11/20

76/76 [=====] - 492s 6s/step - loss: 0.3860 - accuracy: 0.8509

Epoch 12/20

76/76 [=====] - 494s 6s/step - loss: 0.3036 - accuracy: 0.8895

Epoch 13/20

76/76 [=====] - 493s 6s/step - loss: 0.2649 - accuracy: 0.9020

Epoch 14/20

76/76 [=====] - 490s 6s/step - loss: 0.2175 - accuracy: 0.9241

Epoch 15/20

```
76/76 [=====] - 489s 6s/step - loss: 0.1590 - accurac
y: 0.9413
Epoch 16/20
76/76 [=====] - 489s 6s/step - loss: 0.1416 - accurac
y: 0.9494
Epoch 17/20
76/76 [=====] - 489s 6s/step - loss: 0.1038 - accurac
y: 0.9637
Epoch 18/20
76/76 [=====] - 489s 6s/step - loss: 0.1187 - accurac
y: 0.9625
Epoch 19/20
76/76 [=====] - 490s 6s/step - loss: 0.0746 - accurac
y: 0.9728
Epoch 20/20
76/76 [=====] - 489s 6s/step - loss: 0.0760 - accurac
y: 0.9738
```

Out[21]: <keras.callbacks.History at 0x201cb55cfd0>

```
In [22]: test_loss, test_acc = model.evaluate(x_test, y_test)
print("Test Accuracy: ",test_acc*100,"%")
```

```
79/79 [=====] - 51s 642ms/step - loss: 1.5483 - accura
cy: 0.7420
Test Accuracy: 74.20127987861633 %
```

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []: