

```
In [ ]: pip install tensorflow
```

```
In [ ]: pip install sklearn
```

```
In [ ]: pip install matplotlib
```

```
In [ ]: pip install keras
```

```
In [ ]: pip install seaborn
```

```
In [ ]: pip install pydotplus
```

```
In [1]: import matplotlib.pyplot as plt
from PIL import Image
import seaborn as sns
import numpy as np
import pandas as pd
import os
from tensorflow.keras.utils import to_categorical
from glob import glob

from sklearn.model_selection import train_test_split
import keras
from keras.models import Sequential
from keras.layers import Dense, Dropout
import tensorflow as tf
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import ReduceLROnPlateau
from tensorflow.keras import layers
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten

from sklearn.metrics import confusion_matrix
import itertools
import matplotlib.pyplot as plt

import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: df = pd.read_csv('skinCancer/HAM10000_metadata_Increased.csv')
df.head(10)
```

Out[2]:

	lesion_id	image_id	dx	dx_type	age	sex	localization
0	HAM_0000118	ISIC_0027419	bkl	histo	80.0	male	scalp
1	HAM_0000118	ISIC_0025030	bkl	histo	80.0	male	scalp
2	HAM_0002730	ISIC_0026769	bkl	histo	80.0	male	scalp
3	HAM_0002730	ISIC_0025661	bkl	histo	80.0	male	scalp
4	HAM_0001466	ISIC_0031633	bkl	histo	75.0	male	ear
5	HAM_0001466	ISIC_0027850	bkl	histo	75.0	male	ear
6	HAM_0002761	ISIC_0029176	bkl	histo	60.0	male	face
7	HAM_0002761	ISIC_0029068	bkl	histo	60.0	male	face
8	HAM_0005132	ISIC_0025837	bkl	histo	70.0	female	back
9	HAM_0005132	ISIC_0025209	bkl	histo	70.0	female	back

```
In [3]: df.isnull().sum()
```

Out[3]:

lesion_id	0
image_id	0
dx	0
dx_type	0
age	119
sex	0
localization	0

dtype: int64

```
In [4]: df.count()
```

Out[4]:

lesion_id	49517
image_id	49517
dx	49517
dx_type	49517
age	49398
sex	49517
localization	49517

dtype: int64

```
In [5]: df['age'].fillna(int(df['age'].mean()),inplace=True)
df.isnull().sum()
```

```
Out[5]: lesion_id      0
image_id      0
dx            0
dx_type       0
age           0
sex           0
localization   0
dtype: int64
```

```
In [6]: lesion_type_dict = {
        'nv': 'Melanocytic nevi',
        'mel': 'Melanoma',
        'bkl': 'Benign keratosis-like lesions ',
        'bcc': 'Basal cell carcinoma',
        'akiec': 'Actinic keratoses',
        'vasc': 'Vascular lesions',
        'df': 'Dermatofibroma'
    }
base_skin_dir = 'skinCancer'

imageid_path_dict = {os.path.splitext(os.path.basename(x))[0]: x
                     for x in glob(os.path.join(base_skin_dir, '*', '*.jpg'))}
```

```
In [7]: df['path'] = df['image_id'].map(imageid_path_dict.get)
df['cell_type'] = df['dx'].map(lesion_type_dict.get)
df['cell_type_idx'] = pd.Categorical(df['cell_type']).codes
df.head()
```

```
Out[7]:
```

	lesion_id	image_id	dx	dx_type	age	sex	localization
0	HAM_0000118	ISIC_0027419	bkl	histo	80.0	male	scalp skinCancer\HAM10000_images_
1	HAM_0000118	ISIC_0025030	bkl	histo	80.0	male	scalp skinCancer\HAM10000_images_
2	HAM_0002730	ISIC_0026769	bkl	histo	80.0	male	scalp skinCancer\HAM10000_images_
3	HAM_0002730	ISIC_0025661	bkl	histo	80.0	male	scalp skinCancer\HAM10000_images_
4	HAM_0001466	ISIC_0031633	bkl	histo	75.0	male	ear skinCancer\HAM10000_images_

```
In [8]: df['image'] = df['path'].map(lambda x: np.asarray(Image.open(x).resize((125,100))
```

```

In [9]: plt.figure(figsize=(20,10))
plt.subplots_adjust(left=0.125, bottom=1, right=0.9, top=2, hspace=0.2)
plt.subplot(2,4,1)
plt.title("AGE", fontsize=15)
plt.ylabel("Count")
df['age'].value_counts().plot.bar()

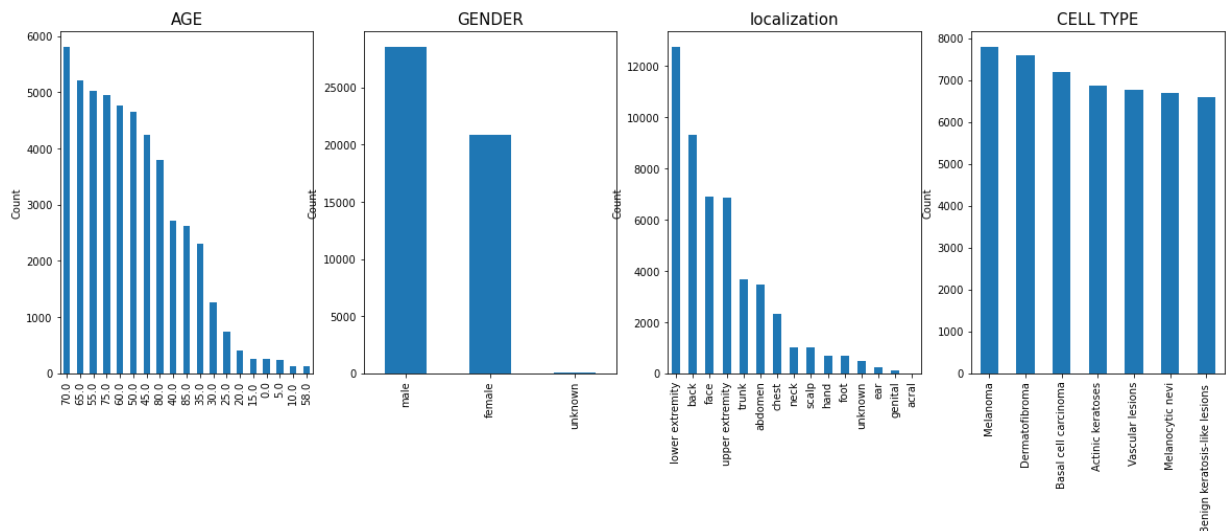
plt.subplot(2,4,2)
plt.title("GENDER", fontsize=15)
plt.ylabel("Count")
df['sex'].value_counts().plot.bar()

plt.subplot(2,4,3)
plt.title("localization", fontsize=15)
plt.ylabel("Count")
plt.xticks(rotation=45)
df['localization'].value_counts().plot.bar()

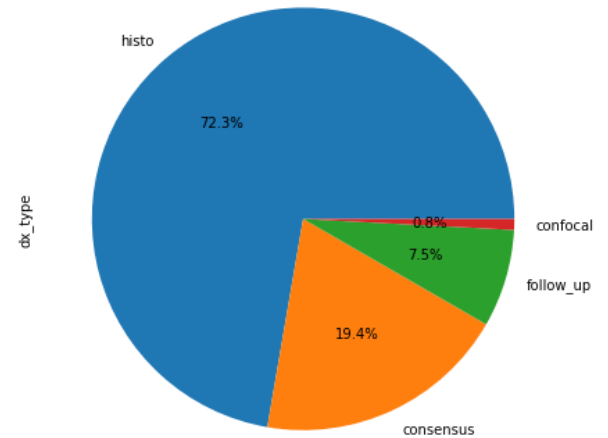
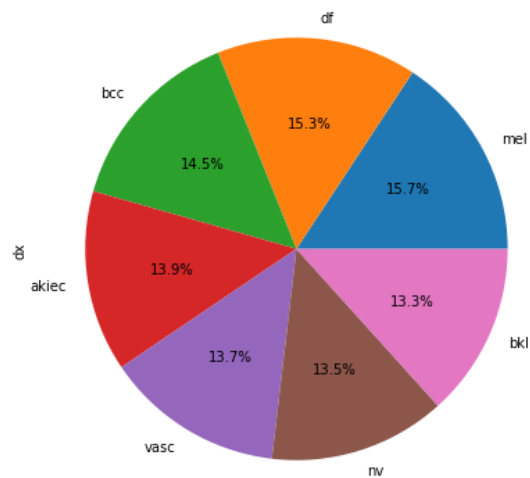
plt.subplot(2,4,4)
plt.title("CELL TYPE", fontsize=15)
plt.ylabel("Count")
df['cell_type'].value_counts().plot.bar()

```

Out[9]: <AxesSubplot:title={'center':'CELL TYPE'}, ylabel='Count'>



```
In [10]: plt.figure(figsize=(15,10))
plt.subplot(1,2,1)
df['dx'].value_counts().plot.pie(autopct="%1.1f%%")
plt.subplot(1,2,2)
df['dx_type'].value_counts().plot.pie(autopct="%1.1f%%")
plt.show()
```



```
In [11]: features=df.drop(columns=['cell_type_idx','image_id'],axis=1)
target=df['cell_type_idx']
features.head()
```

Out[11]:

	lesion_id	dx	dx_type	age	sex	localization	
0	HAM_0000118	bkl	histo	80.0	male	scalp	skinCancer\HAM10000_images_part_1\ISIC_002
1	HAM_0000118	bkl	histo	80.0	male	scalp	skinCancer\HAM10000_images_part_1\ISIC_002
2	HAM_0002730	bkl	histo	80.0	male	scalp	skinCancer\HAM10000_images_part_1\ISIC_002
3	HAM_0002730	bkl	histo	80.0	male	scalp	skinCancer\HAM10000_images_part_1\ISIC_002
4	HAM_0001466	bkl	histo	75.0	male	ear	skinCancer\HAM10000_images_part_2\ISIC_003



```

In [12]: x_train_o, x_test_o, y_train_o, y_test_o = train_test_split(features, target, test_size=0.2,
                             random_state=42, stratify=tf.unique(x_train_o.cell_type.values))

Out[12]: Unique(y=<tf.Tensor: shape=(7,), dtype=string, numpy=
array([b'Actinic keratoses', b'Basal cell carcinoma', b'Dermatofibroma',
      b'Melanoma', b'Benign keratosis-like lesions ',
      b'Melanocytic nevi', b'Vascular lesions'], dtype=object)>, idx=<tf.Tensor: shape=(37137,), dtype=int32, numpy=array([0, 1, 1, ..., 1, 1, 6])>)

In [13]: x_train = np.asarray(x_train_o['image'].tolist())
x_test = np.asarray(x_test_o['image'].tolist())

x_train_mean = np.mean(x_train)
x_train_std = np.std(x_train)

x_test_mean = np.mean(x_test)
x_test_std = np.std(x_test)

x_train = (x_train - x_train_mean)/x_train_std
x_test = (x_test - x_test_mean)/x_test_std

In [14]: # Perform one-hot encoding on the labels
y_train = to_categorical(y_train_o, num_classes = 7)
y_test = to_categorical(y_test_o, num_classes = 7)
y_test

Out[14]: array([[0., 0., 0., ..., 0., 1., 0.],
                [0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 1.],
                ...,
                [0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 1.],
                [0., 0., 0., ..., 0., 0., 0.]], dtype=float32)

In [15]: y_test[1]

Out[15]: array([0., 0., 0., 1., 0., 0., 0.], dtype=float32)

In [16]: y_train

Out[16]: array([[1., 0., 0., ..., 0., 0., 0.],
                [0., 1., 0., ..., 0., 0., 0.],
                [0., 1., 0., ..., 0., 0., 0.],
                ...,
                [0., 1., 0., ..., 0., 0., 0.],
                [0., 1., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 1.]], dtype=float32)

```

```
In [17]: y_test_o.value_counts()
```

```
Out[17]: 5    1924
         3    1907
         1    1750
         0    1730
         6    1697
         2    1689
         4    1683
         Name: cell_type_idx, dtype: int64
```

```
In [18]: # x_train, x_validate, y_train, y_validate = train_test_split(x_train, y_train, t

# Reshape image in 3 dimensions (height = 100, width = 125 , canal = 3)
x_train = x_train.reshape(x_train_o.shape[0], *(100, 125, 3))
x_test = x_test.reshape(x_test_o.shape[0], *(100, 125, 3))
```

```
In [19]: x_train = x_train.reshape(x_train.shape[0],125*100*3)
         x_test = x_test.reshape(x_test.shape[0],125*100*3)
         print(x_train.shape)
         print(x_test.shape)
```

```
(37137, 37500)
(12380, 37500)
```

```
In [20]: print(x_train)
```

```
[[ 0.67592044 -0.3525907  0.07794885 ...  0.96294681  0.12578657
  0.31713748]
 [-0.04164547 -0.85488684 -1.54853389 ...  0.1736243 -0.6874548
 -1.14191321]
 [ 0.07794885 -1.33326412 -1.45285844 ...  0.60416385 -0.23299638
 -0.16123979]
 ...
 [-2.93582799 -3.31852981 -2.79231481 ... -3.17501663 -3.41420527
 -3.03150345]
 [ 0.53240726 -0.63961707 -0.71137366 ...  0.19754317 -1.02231889
 -1.14191321]
 [ 0.93902794 -0.30475297  0.19754317 ...  0.65200158 -0.63961707
 -0.23299638]]
```

Decision Tree


```

In [24]: depth = range(1,51,5)
testing_accuracy = []
training_accuracy = []
score = 0

for i in depth:
    tree = DecisionTreeClassifier(max_depth = i, criterion = 'entropy')
    tree.fit(x_train, y_train)

    y_predict_train = tree.predict(x_train)
    training_accuracy.append(accuracy_score(y_train, y_predict_train))

    y_predict_test = tree.predict(x_test)
    acc_score = accuracy_score(y_test,y_predict_test)
    testing_accuracy.append(acc_score)

    print(i)

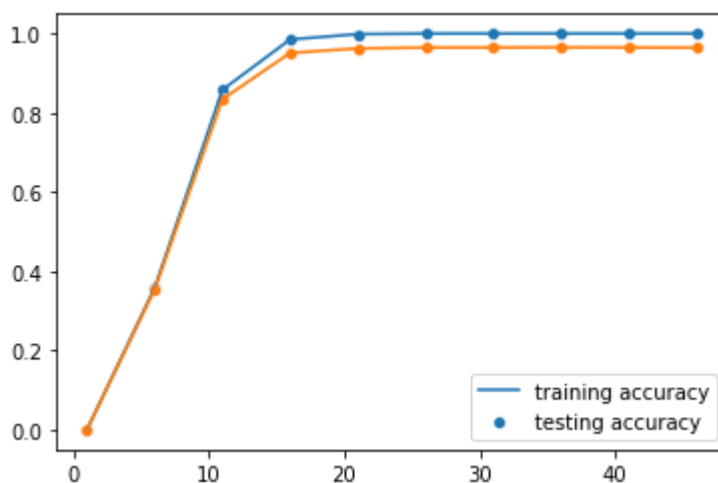
    if score < acc_score:
        score = acc_score
        best_depth = i

sns.lineplot(depth, training_accuracy)
sns.scatterplot(depth, training_accuracy)
sns.lineplot(depth, testing_accuracy)
sns.scatterplot(depth, testing_accuracy)
plt.legend(['training accuracy', 'testing accuracy'])

```

1
6
11
16
21
26
31
36
41
46

Out[24]: <matplotlib.legend.Legend at 0x14e31eae90>



```
In [25]: print('This is the best depth for Decision Tree Classifier: ', best_depth, '\nAcc
result = confusion_matrix(y_test.argmax(axis=1), y_predict_test.argmax(axis=1),)
print("Confusion Matrix:", "\n", result)
report = classification_report(y_test.argmax(axis=1), y_predict_test.argmax(axis=1),)
print("Classification Report:" , "\n", report)
```

This is the best depth for Decision Tree Classifier: 36

Accuracy score is: 0.9651050080775444

Confusion Matrix:

```
[[1702    0    7    0   21    0    0]
 [    0 1736    0    0   11    0    3]
 [    0    0 1689    0    0    0    0]
 [    0    0    0 1907    0    0    0]
 [   22   41  126    7 1284  176   27]
 [    0    0    2    0    0 1922    0]
 [    0    0    0    0    0    0 1697]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.99	0.98	0.99	1730
1	0.98	0.99	0.98	1750
2	0.93	1.00	0.96	1689
3	1.00	1.00	1.00	1907
4	0.98	0.76	0.86	1683
5	0.92	1.00	0.96	1924
6	0.98	1.00	0.99	1697
accuracy			0.96	12380
macro avg	0.97	0.96	0.96	12380
weighted avg	0.97	0.96	0.96	12380

Random Forest

```

In [26]: #random Forest
depth = range(5,51,5)
testing_accuracy = []
training_accuracy = []
score = 0

for i in depth:
    tree = RandomForestClassifier(max_depth = i, criterion = 'gini', random_state=42)
    tree.fit(x_train, y_train)

    y_predict_train = tree.predict(x_train)
    training_accuracy.append(accuracy_score(y_train, y_predict_train))

    y_predict_test = tree.predict(x_test)
    acc_score = accuracy_score(y_test,y_predict_test)
    testing_accuracy.append(acc_score)

    print(i)

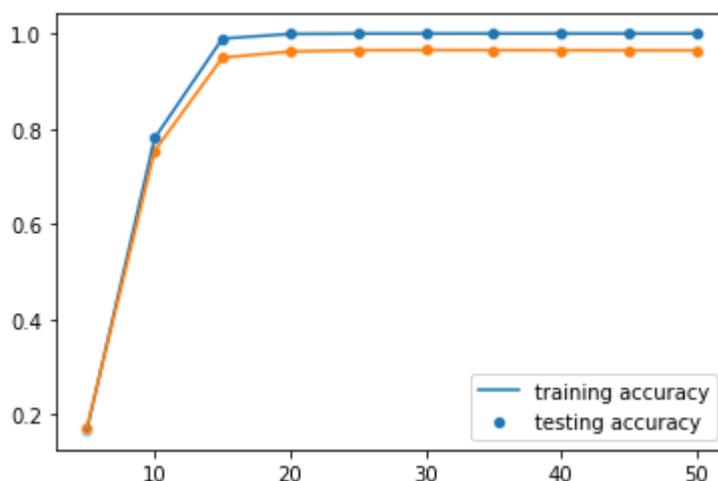
    if score < acc_score:
        score = acc_score
        best_depth = i

sns.lineplot(depth, training_accuracy)
sns.scatterplot(depth, training_accuracy)
sns.lineplot(depth, testing_accuracy)
sns.scatterplot(depth, testing_accuracy)
plt.legend(['training accuracy', 'testing accuracy'])

```

5
10
15
20
25
30
35
40
45
50

Out[26]: <matplotlib.legend.Legend at 0x14e31f2f880>



```
In [27]: print('This is the best depth for Decision Tree Classifier: ', best_depth, '\nAcc
result = confusion_matrix(y_test.argmax(axis=1), y_predict_test.argmax(axis=1),)
print("Confusion Matrix:", "\n", result)
report = classification_report(y_test.argmax(axis=1), y_predict_test.argmax(axis=1),)
print("Classification Report:" , "\n", report)
```

This is the best depth for Decision Tree Classifier: 30

Accuracy score is: 0.9652665589660743

Confusion Matrix:

```
[[1730    0    0    0    0    0    0]
 [   0 1750    0    0    0    0    0]
 [   0    0 1689    0    0    0    0]
 [   0    0    0 1907    0    0    0]
 [ 398    3   15    0 1245   22    0]
 [   0    0    0    0    0 1924    0]
 [   0    0    0    0    0    0 1697]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.81	1.00	0.90	1730
1	1.00	1.00	1.00	1750
2	0.99	1.00	1.00	1689
3	1.00	1.00	1.00	1907
4	1.00	0.74	0.85	1683
5	0.99	1.00	0.99	1924
6	1.00	1.00	1.00	1697
accuracy			0.96	12380
macro avg	0.97	0.96	0.96	12380
weighted avg	0.97	0.96	0.96	12380

KNN

```
In [21]: k = range(1,102,3)
testing_accuracy = []
training_accuracy = []
score = 0

for i in k:
    knn = KNeighborsClassifier(n_neighbors = i)
    knn.fit(x_train, y_train)

    y_predict_train = knn.predict(x_train)
    training_accuracy.append(accuracy_score(y_train, y_predict_train))

    y_predict_test = knn.predict(x_test)
    acc_score = accuracy_score(y_test,y_predict_test)
    testing_accuracy.append(acc_score)

    print(i)

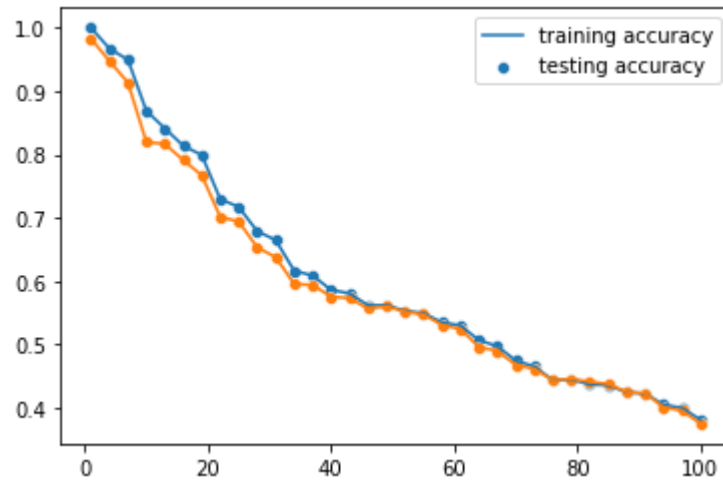
    if score < acc_score:
        score = acc_score
        best_k = i

sns.lineplot(k, training_accuracy)
sns.scatterplot(k, training_accuracy)
sns.lineplot(k, testing_accuracy)
sns.scatterplot(k, testing_accuracy)
plt.legend(['training accuracy', 'testing accuracy'])
```

1
4
7
10
13
16
19
22
25
28
31
34
37
40
43
46
49
52
55
58
61
64
67
70
73
76
79
82

85
88
91
94
97
100

Out[21]: <matplotlib.legend.Legend at 0x12a87a9ac10>



```
In [28]: '''print('This is the best depth for Decision Tree Classifier: ', best_depth, '\n')
result = confusion_matrix(y_test.argmax(axis=1), y_predict_test.argmax(axis=1),)
print("Confusion Matrix:", "\n", result)
report = classification_report(y_test.argmax(axis=1), y_predict_test.argmax(axis=1),)
print("Classification Report:" , "\n", report)'''
```

Confusion Matrix:

```
[[1259  114   35  237    6    0   79]
 [ 862  476   45  279    0    0   88]
 [ 906   69  345  237   29   53   50]
 [ 543   27    0 1337    0    0    0]
 [ 627   25   34  158  686   58   95]
[1248    6  137  105   45  289   94]
 [ 342   12    0   57   29    0 1257]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.22	0.73	0.33	1730
1	0.65	0.27	0.38	1750
2	0.58	0.20	0.30	1689
3	0.55	0.70	0.62	1907
4	0.86	0.41	0.55	1683
5	0.72	0.15	0.25	1924
6	0.76	0.74	0.75	1697
accuracy			0.46	12380
macro avg	0.62	0.46	0.46	12380
weighted avg	0.62	0.46	0.45	12380

```
In [24]: x_train.shape
```

```
Out[24]: (37137, 37500)
```

MLP

In [25]: *# define the keras model*

```

model = Sequential()

model.add(Dense(units= 64, kernel_initializer = 'uniform', activation = 'relu', i
model.add(Dense(units= 128, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dense(units= 64, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dense(units= 32, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dense(units = 7, kernel_initializer = 'uniform', activation = 'softmax'
model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 64)	2400064
dense_1 (Dense)	(None, 128)	8320
dense_2 (Dense)	(None, 64)	8256
dense_3 (Dense)	(None, 32)	2080
dense_4 (Dense)	(None, 7)	231

```

=====
Total params: 2,418,951
Trainable params: 2,418,951
Non-trainable params: 0
=====

```



```
In [26]: # compile the keras model
model.compile(optimizer = "Adam", loss = 'categorical_crossentropy', metrics = ['accuracy'])

# fit the keras model on the dataset
history = model.fit(x_train, y_train, batch_size = 20, epochs = 25)
```

Epoch 1/25
1857/1857 [=====] - 26s 13ms/step - loss: 1.1413 - accuracy: 0.5545

Epoch 2/25
1857/1857 [=====] - 25s 14ms/step - loss: 0.6530 - accuracy: 0.7451

Epoch 3/25
1857/1857 [=====] - 25s 14ms/step - loss: 0.4905 - accuracy: 0.8061

Epoch 4/25
1857/1857 [=====] - 25s 14ms/step - loss: 0.4207 - accuracy: 0.8389

Epoch 5/25
1857/1857 [=====] - 23s 12ms/step - loss: 0.3594 - accuracy: 0.8637

Epoch 6/25
1857/1857 [=====] - 23s 12ms/step - loss: 0.3272 - accuracy: 0.8782

Epoch 7/25
1857/1857 [=====] - 25s 13ms/step - loss: 0.2870 - accuracy: 0.8936

Epoch 8/25
1857/1857 [=====] - 25s 13ms/step - loss: 0.2702 - accuracy: 0.9006

Epoch 9/25
1857/1857 [=====] - 25s 13ms/step - loss: 0.2527 - accuracy: 0.9080

Epoch 10/25
1857/1857 [=====] - 25s 13ms/step - loss: 0.2223 - accuracy: 0.9213

Epoch 11/25
1857/1857 [=====] - 25s 14ms/step - loss: 0.2130 - accuracy: 0.9216

Epoch 12/25
1857/1857 [=====] - 25s 13ms/step - loss: 0.2102 - accuracy: 0.9269

Epoch 13/25
1857/1857 [=====] - 25s 13ms/step - loss: 0.1950 - accuracy: 0.9327

Epoch 14/25
1857/1857 [=====] - 25s 14ms/step - loss: 0.1774 - accuracy: 0.9366

Epoch 15/25
1857/1857 [=====] - 26s 14ms/step - loss: 0.1805 - accuracy: 0.9371

Epoch 16/25
1857/1857 [=====] - 26s 14ms/step - loss: 0.1659 - accuracy: 0.9420

Epoch 17/25
1857/1857 [=====] - 26s 14ms/step - loss: 0.1665 - accuracy: 0.9437

```

Epoch 18/25
1857/1857 [=====] - 26s 14ms/step - loss: 0.1453 - a
ccuracy: 0.9501
Epoch 19/25
1857/1857 [=====] - 27s 14ms/step - loss: 0.1458 - a
ccuracy: 0.9507
Epoch 20/25
1857/1857 [=====] - 26s 14ms/step - loss: 0.1470 - a
ccuracy: 0.9504
Epoch 21/25
1857/1857 [=====] - 26s 14ms/step - loss: 0.1352 - a
ccuracy: 0.9552
Epoch 22/25
1857/1857 [=====] - 26s 14ms/step - loss: 0.1305 - a
ccuracy: 0.9572
Epoch 23/25
1857/1857 [=====] - 26s 14ms/step - loss: 0.1333 - a
ccuracy: 0.9554
Epoch 24/25
1857/1857 [=====] - 27s 14ms/step - loss: 0.1134 - a
ccuracy: 0.9616
Epoch 25/25
1857/1857 [=====] - 26s 14ms/step - loss: 0.1235 - a
ccuracy: 0.9584

```

```

In [27]: accuracy = model.evaluate(x_test, y_test, verbose=1)[1]
print("Test: accuracy = ",accuracy*100,"%")

387/387 [=====] - 5s 12ms/step - loss: 0.1793 - accura
cy: 0.9462
Test: accuracy = 94.62035298347473 %

```

CNN

```

In [ ]: x_train = x_train.reshape(x_train.shape[0], 125,100,3)
x_test = x_test.reshape(x_test.shape[0], 125, 100, 3)
print(x_train.shape)
print(x_test.shape)

```

```
In [ ]: input_shape = (125, 100, 3)
        num_classes = 7

        model = Sequential()
        model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', padding = 'Same', input_shape=input_shape))
        model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', padding = 'Same'))
        model.add(MaxPooling2D(pool_size = (2, 2)))
        model.add(Dropout(0.16))

        model.add(Conv2D(64, (3, 3), activation='relu', padding = 'same'))
        model.add(Conv2D(64, (3, 3), activation='relu', padding = 'Same'))
        model.add(MaxPooling2D(pool_size=(2, 2)))
        model.add(Dropout(0.20))

        model.add(Flatten())
        model.add(Dense(256, activation='relu'))
        model.add(Dense(128, activation='relu'))
        model.add(Dropout(0.3))
        model.add(Dense(num_classes, activation='softmax'))
        model.summary()
```

```
In [ ]: model.compile(optimizer= "Adam",
                      loss='categorical_crossentropy',
                      metrics=['accuracy'])

        learning_rate_reduction = ReduceLRonPlateau(monitor='val_accuracy',
                                                    patience=4,
                                                    verbose=1,
                                                    factor=0.5,
                                                    min_lr=0.00001)
```

```
In [ ]: #x_train, x_validate, y_train, y_validate = train_test_split(x_train, y_train, test_size=0.2, random_state=42)

        # Reshape image in 3 dimensions (height = 100, width = 125 , canal = 3)
        x_train = x_train.reshape(x_train.shape[0], *(100, 125, 3))
        x_test = x_test.reshape(x_test.shape[0], *(100, 125, 3))
        #x_validate = x_validate.reshape(x_validate.shape[0], *(100, 125, 3))

        model.fit(x_train, y_train, batch_size = 100, epochs = 20)
```

```
In [ ]: test_loss, test_acc = model.evaluate(x_test, y_test)
        print("Test Accuracy: ", test_acc*100, "%")
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

In []:

In []:

In []:

In []:

In []: