

In [2]: `pip install seaborn`

```
Requirement already satisfied: seaborn in c:\programdata\anaconda3\lib\site-packages (0.11.2)
Requirement already satisfied: matplotlib>=2.2 in c:\programdata\anaconda3\lib\site-packages (from seaborn) (3.5.1)
Requirement already satisfied: numpy>=1.15 in c:\programdata\anaconda3\lib\site-packages (from seaborn) (1.21.5)
Requirement already satisfied: scipy>=1.0 in c:\programdata\anaconda3\lib\site-packages (from seaborn) (1.7.3)
Requirement already satisfied: pandas>=0.23 in c:\programdata\anaconda3\lib\site-packages (from seaborn) (1.4.2)
Requirement already satisfied: pillow>=6.2.0 in c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.2->seaborn) (9.0.1)
Requirement already satisfied: packaging>=20.0 in c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.2->seaborn) (21.3)
Requirement already satisfied: pyparsing>=2.2.1 in c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.2->seaborn) (3.0.4)
Requirement already satisfied: python-dateutil>=2.7 in c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.2->seaborn) (2.8.2)
Requirement already satisfied: fonttools>=4.22.0 in c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.2->seaborn) (4.25.0)
Requirement already satisfied: cycler>=0.10 in c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.2->seaborn) (0.11.0)
Requirement already satisfied: kiwisolver>=1.0.1 in c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.2->seaborn) (1.3.2)
Requirement already satisfied: pytz>=2020.1 in c:\programdata\anaconda3\lib\site-packages (from pandas>=0.23->seaborn) (2021.3)
Requirement already satisfied: six>=1.5 in c:\programdata\anaconda3\lib\site-packages (from python-dateutil>=2.7->matplotlib>=2.2->seaborn) (1.16.0)
Note: you may need to restart the kernel to use updated packages.
```

In [3]: `pip install pydotplus`

```
Requirement already satisfied: pydotplus in c:\programdata\anaconda3\lib\site-packages (2.0.2)
Note: you may need to restart the kernel to use updated packages.
Requirement already satisfied: pyparsing>=2.0.1 in c:\programdata\anaconda3\lib\site-packages (from pydotplus) (3.0.4)
```

```
In [1]: import matplotlib.pyplot as plt
from PIL import Image
import seaborn as sns
import numpy as np
import pandas as pd
import os
from tensorflow.keras.utils import to_categorical
from glob import glob

from sklearn.model_selection import train_test_split
import keras
from keras.models import Sequential
from keras.layers import Dense, Dropout
import tensorflow as tf
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import ReduceLROnPlateau
from tensorflow.keras import layers
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten

from sklearn.metrics import confusion_matrix
import itertools
import matplotlib.pyplot as plt

import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: df = pd.read_csv('skinCancer/abstract_metadata.csv')
df.head(10)
```

Out[2]:

	lesion_id	image_id	dx	dx_type	age	sex	localization
0	HAM_0000673	ISIC_0029659	akiec	histo	70	female	face
1	HAM_0005282	ISIC_0025178	akiec	histo	65	male	lower extremity
2	HAM_0006002	ISIC_0029915	akiec	histo	50	female	face
3	HAM_0000549	ISIC_0029360	akiec	histo	70	male	upper extremity
4	HAM_0000549	ISIC_0026152	akiec	histo	70	male	upper extremity
5	HAM_0006875	ISIC_0026575	akiec	histo	80	male	face
6	HAM_0006875	ISIC_0030586	akiec	histo	80	male	face
7	HAM_0002644	ISIC_0029417	akiec	histo	80	female	neck
8	HAM_0005282	ISIC_0028730	akiec	histo	65	male	lower extremity
9	HAM_0006898	ISIC_0029041	akiec	histo	80	male	scalp

```
In [3]: df.isnull().sum()
```

```
Out[3]: lesion_id      0
image_id      0
dx            0
dx_type       0
age           0
sex           0
localization  0
dtype: int64
```

```
In [4]: df.count()
```

```
Out[4]: lesion_id      700
image_id      700
dx            700
dx_type       700
age           700
sex           700
localization  700
dtype: int64
```

```
In [5]: df['age'].fillna(int(df['age'].mean()),inplace=True)
df.isnull().sum()
```

```
Out[5]: lesion_id      0
image_id      0
dx            0
dx_type       0
age           0
sex           0
localization  0
dtype: int64
```

```
In [6]: lesion_type_dict = {
    'nv': 'Melanocytic nevi',
    'mel': 'Melanoma',
    'bkl': 'Benign keratosis-like lesions ',
    'bcc': 'Basal cell carcinoma',
    'akiec': 'Actinic keratoses',
    'vasc': 'Vascular lesions',
    'df': 'Dermatofibroma'
}
base_skin_dir = 'skinCancer'

imageid_path_dict = {os.path.splitext(os.path.basename(x))[0]: x
                      for x in glob(os.path.join(base_skin_dir, '*', '*.jpg'))}
```

```
In [7]: df['path'] = df['image_id'].map(imageid_path_dict.get)
df['cell_type'] = df['dx'].map(lesion_type_dict.get)
df['cell_type_idx'] = pd.Categorical(df['cell_type']).codes
df.head()
```

Out[7]:

	lesion_id	image_id	dx	dx_type	age	sex	localization	
0	HAM_0000673	ISIC_0029659	akiec	histo	70	female	face	skinCancer\HAM10000_image0000673.jpg
1	HAM_0005282	ISIC_0025178	akiec	histo	65	male	lower extremity	skinCancer\HAM10000_image0005282.jpg
2	HAM_0006002	ISIC_0029915	akiec	histo	50	female	face	skinCancer\HAM10000_image0006002.jpg
3	HAM_0000549	ISIC_0029360	akiec	histo	70	male	upper extremity	skinCancer\HAM10000_image0000549.jpg
4	HAM_0000549	ISIC_0026152	akiec	histo	70	male	upper extremity	skinCancer\HAM10000_image0000549.jpg

In [8]: df['image'] = df['path'].map(lambda x: np.asarray(Image.open(x).resize((125,100))))

```

In [9]: plt.figure(figsize=(20,10))
plt.subplots_adjust(left=0.125, bottom=1, right=0.9, top=2, hspace=0.2)
plt.subplot(2,4,1)
plt.title("AGE",fontsize=15)
plt.ylabel("Count")
df['age'].value_counts().plot.bar()


plt.subplot(2,4,2)
plt.title("GENDER",fontsize=15)
plt.ylabel("Count")
df['sex'].value_counts().plot.bar()

plt.subplot(2,4,3)
plt.title("localization",fontsize=15)
plt.ylabel("Count")
plt.xticks(rotation=45)
df['localization'].value_counts().plot.bar()



plt.subplot(2,4,4)
plt.title("CELL TYPE",fontsize=15)
plt.ylabel("Count")
df['cell_type'].value_counts().plot.bar()

```

Out[9]: <AxesSubplot:title={'center':'CELL TYPE'}, ylabel='Count'>



```
In [10]: plt.figure(figsize=(15,10))
plt.subplot(1,2,1)
df['dx'].value_counts().plot.pie(autopct="%1.1f%%")
plt.subplot(1,2,2)
df['dx_type'].value_counts().plot.pie(autopct="%1.1f%%")
plt.show()
```



```
In [11]: features=df.drop(columns=['cell_type_idx','image_id'],axis=1)
target=df['cell_type_idx']
features.head()
```

Out[11]:

	lesion_id	dx	dx_type	age	sex	localization	
0	HAM_0000673	akiec	histo	70	female	face	skinCancer\HAM10000_images_part_2\ISIC_
1	HAM_0005282	akiec	histo	65	male	lower extremity	skinCancer\HAM10000_images_part_1\ISIC_
2	HAM_0006002	akiec	histo	50	female	face	skinCancer\HAM10000_images_part_2\ISIC_
3	HAM_0000549	akiec	histo	70	male	upper extremity	skinCancer\HAM10000_images_part_2\ISIC_
4	HAM_0000549	akiec	histo	70	male	upper extremity	skinCancer\HAM10000_images_part_1\ISIC_

```
In [12]: x_train_o, x_test_o, y_train_o, y_test_o = train_test_split(features, target, test_size=0.2,
tf.unique(x_train_o.cell_type.values)
```

```
Out[12]: Unique(y=<tf.Tensor: shape=(7,), dtype=string, numpy=
array([b'Melanoma', b'Benign keratosis-like lesions ',
      b'Melanocytic nevi', b'Actinic keratoses', b'Dermatofibroma',
      b'Basal cell carcinoma', b'Vascular lesions']), dtype=object)>, idx=<tf.Tensor: shape=(525,), dtype=int32, numpy=
array([0, 1, 0, 0, 0, 2, 3, 4, 0, 0, 1, 0, 3, 1, 2, 3, 5, 3, 2, 1, 4, 3,
      3, 1, 6, 0, 1, 4, 4, 0, 2, 6, 6, 0, 3, 5, 4, 2, 4, 6, 6, 2, 6, 3,
      1, 6, 2, 4, 3, 2, 5, 4, 1, 5, 1, 6, 2, 5, 3, 2, 5, 0, 0, 0, 4, 6,
      4, 4, 1, 4, 0, 2, 0, 0, 3, 5, 6, 6, 3, 5, 3, 3, 0, 0, 1, 1, 1, 6,
      5, 2, 3, 5, 3, 1, 4, 3, 0, 1, 5, 5, 0, 0, 6, 5, 3, 4, 0, 4, 4, 6,
      6, 4, 3, 4, 0, 5, 6, 1, 2, 5, 2, 3, 5, 4, 5, 0, 0, 2, 2, 3, 0, 6,
      4, 2, 3, 2, 5, 2, 1, 1, 6, 3, 0, 5, 0, 5, 2, 6, 6, 2, 6, 2, 2, 6,
      6, 1, 0, 2, 4, 2, 6, 1, 6, 4, 6, 2, 5, 1, 3, 2, 3, 3, 1, 5, 3, 3,
      4, 0, 2, 3, 5, 5, 1, 3, 3, 5, 4, 6, 0, 6, 3, 5, 6, 4, 0, 2, 2, 1,
      2, 6, 2, 4, 3, 1, 1, 5, 2, 0, 5, 6, 3, 4, 0, 5, 5, 3, 5, 2, 1, 3,
      6, 6, 2, 1, 6, 0, 1, 1, 4, 2, 6, 0, 6, 4, 1, 1, 5, 6, 6, 1, 0, 0,
      6, 4, 4, 4, 2, 1, 2, 2, 1, 6, 4, 0, 4, 3, 0, 1, 6, 5, 2, 5, 6, 0,
      1, 5, 6, 4, 3, 6, 6, 1, 2, 6, 5, 0, 4, 0, 2, 3, 3, 4, 6, 4, 6, 3,
      0, 4, 1, 0, 5, 0, 2, 5, 5, 2, 4, 2, 5, 2, 0, 6, 2, 0, 5, 1, 2, 5,
      1, 6, 2, 0, 3, 1, 4, 1, 2, 0, 4, 0, 6, 5, 6, 1, 1, 2, 0, 0, 5, 0,
      2, 3, 2, 1, 2, 0, 1, 2, 4, 5, 1, 4, 1, 2, 5, 2, 0, 4, 2, 2, 5, 4,
      5, 1, 4, 2, 1, 4, 1, 3, 0, 5, 6, 3, 4, 4, 6, 3, 3, 3, 4, 1, 4, 2,
      5, 5, 0, 6, 6, 0, 6, 4, 4, 2, 5, 1, 1, 4, 0, 5, 4, 3, 1, 1, 3, 5,
      6, 3, 6, 1, 0, 4, 5, 5, 2, 3, 3, 6, 0, 5, 2, 4, 6, 0, 1, 5, 3, 2,
      0, 2, 5, 0, 3, 1, 5, 4, 4, 6, 0, 3, 0, 6, 5, 1, 1, 5, 0, 3, 2, 5,
      0, 6, 0, 6, 6, 1, 5, 1, 1, 5, 4, 0, 5, 1, 0, 0, 5, 4, 3, 2, 6, 5,
      6, 4, 3, 3, 6, 5, 2, 4, 3, 3, 0, 4, 2, 3, 3, 2, 3, 5, 3, 0, 4, 2,
      6, 3, 6, 0, 3, 1, 2, 2, 4, 5, 2, 0, 6, 5, 3, 4, 4, 1, 2, 6, 4, 3,
      4, 6, 4, 6, 3, 1, 3, 1, 0, 4, 3, 0, 6, 3, 1, 1, 3, 1, 6]))>)
```

```
In [13]: x_train = np.asarray(x_train_o['image'].tolist())
x_test = np.asarray(x_test_o['image'].tolist())

x_train_mean = np.mean(x_train)
x_train_std = np.std(x_train)

x_test_mean = np.mean(x_test)
x_test_std = np.std(x_test)

x_train = (x_train - x_train_mean)/x_train_std
x_test = (x_test - x_test_mean)/x_test_std
```



```
In [14]: # Perform one-hot encoding on the labels
y_train = to_categorical(y_train_o, num_classes = 7)
y_test = to_categorical(y_test_o, num_classes = 7)
y_test
```

```
Out[14]: array([[0., 1., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 1.],
                [1., 0., 0., ..., 0., 0., 0.],
                ...,
                [0., 0., 0., ..., 1., 0., 0.],
                [0., 0., 0., ..., 0., 0., 1.],
                [0., 0., 0., ..., 0., 0., 1.]], dtype=float32)
```

```
In [15]: y_test[1]
```

```
Out[15]: array([0., 0., 0., 0., 0., 0., 1.], dtype=float32)
```

```
In [16]: y_train
```

```
Out[16]: array([[0., 0., 0., ..., 0., 1., 0.],
                [0., 0., 1., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 1., 0.],
                ...,
                [1., 0., 0., ..., 0., 0., 0.],
                [0., 0., 1., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 1.]], dtype=float32)
```

```
In [17]: y_test_o.value_counts()
```

```
Out[17]: 1    28
         2    27
         3    27
         0    25
         4    24
         6    23
         5    21
         Name: cell_type_idx, dtype: int64
```

```
In [18]: # x_train, x_validate, y_train, y_validate = train_test_split(x_train, y_train, t

# Reshape image in 3 dimensions (height = 100, width = 125 , canal = 3)
x_train = x_train.reshape(x_train_o.shape[0], *(100, 125, 3))
x_test = x_test.reshape(x_test_o.shape[0], *(100, 125, 3))
```

```
In [19]: x_train = x_train.reshape(x_train.shape[0],125*100*3)
x_test = x_test.reshape(x_test.shape[0],125*100*3)
print(x_train.shape)
print(x_test.shape)
```

```
(525, 37500)
(175, 37500)
```

```
In [20]: print(x_train)
```

```
[[-0.09140732 -0.49162806 -0.46808566 ...  0.23818623 -0.02078013
  0.12047425]
 [ 0.07338946 -0.79767921 -0.09140732 ...  0.04984706 -0.44454326
  0.16755904]
 [-0.11494971 -0.44454326 -0.2091193  ...  0.02630467 -0.42100087
 -0.1855769 ]
 ...
 [ 1.22696689  0.23818623  0.70903416 ...  1.58010283 -0.1855769
  0.54423739]
 [-0.13849211 -0.70350962 -0.65642483 ... -0.32683128 -0.72705202
 -0.75059442]
 [ 1.60364523  0.61486458  0.70903416 ...  1.55656044  0.77966135
  0.77966135]]
```

Decision Tree

```

In [21]: depth = range(5,51,5)
testing_accuracy = []
training_accuracy = []
score = 0

for i in depth:
    tree = DecisionTreeClassifier(max_depth = i, criterion = 'entropy')
    tree.fit(x_train, y_train)

    y_predict_train = tree.predict(x_train)
    training_accuracy.append(accuracy_score(y_train, y_predict_train))

    y_predict_test = tree.predict(x_test)
    acc_score = accuracy_score(y_test, y_predict_test)
    testing_accuracy.append(acc_score)

    print(i)


    if score < acc_score:
        score = acc_score
        best_depth = i

sns.lineplot(depth, training_accuracy)
sns.scatterplot(depth, training_accuracy)
sns.lineplot(depth, testing_accuracy)
sns.scatterplot(depth, testing_accuracy)
plt.legend(['training accuracy', 'testing accuracy'])

```

5
10
15
20
25
30
35
40
45
50

Out[21]: <matplotlib.legend.Legend at 0x296633e2040>



```
In [22]: print('This is the best depth for Decision Tree Classifier: ', best_depth, '\nAcc
result = confusion_matrix(y_test.argmax(axis=1), y_predict_test.argmax(axis=1),)
print("Confusion Matrix:", "\n", result)
report = classification_report(y_test.argmax(axis=1), y_predict_test.argmax(axis=
print("Classification Report:" , "\n", report)
```

This is the best depth for Decision Tree Classifier: 50

Accuracy score is: 0.42857142857142855

Confusion Matrix:

```
[[13  4  0  2  4  1  1]
 [ 5  7  2  8  3  1  2]
 [ 1  0 16  1  1  4  4]
 [ 6  2  2 12  0  2  3]
 [ 3  4  1  3 10  2  1]
 [ 5  0  4  1  1  8  2]
 [ 1  1  3  2  5  2  9]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.38	0.52	0.44	25
1	0.39	0.25	0.30	28
2	0.57	0.59	0.58	27
3	0.41	0.44	0.43	27
4	0.42	0.42	0.42	24
5	0.40	0.38	0.39	21
6	0.41	0.39	0.40	23
accuracy			0.43	175
macro avg	0.43	0.43	0.42	175
weighted avg	0.43	0.43	0.42	175

Random Forest

```

In [23]: #random Forest
depth = range(5,51,5)
testing_accuracy = []
training_accuracy = []
score = 0

for i in depth:
    tree = RandomForestClassifier(max_depth = i, criterion = 'gini', random_state=42)
    tree.fit(x_train, y_train)

    y_predict_train = tree.predict(x_train)
    training_accuracy.append(accuracy_score(y_train, y_predict_train))

    y_predict_test = tree.predict(x_test)
    acc_score = accuracy_score(y_test,y_predict_test)
    testing_accuracy.append(acc_score)

    print(i)


    if score < acc_score:
        score = acc_score
        best_depth = i

sns.lineplot(depth, training_accuracy)
sns.scatterplot(depth, training_accuracy)
sns.lineplot(depth, testing_accuracy)
sns.scatterplot(depth, testing_accuracy)
plt.legend(['training accuracy', 'testing accuracy'])

```

5
10
15
20
25
30
35
40
45
50

Out[23]: <matplotlib.legend.Legend at 0x296634e1af0>



```
In [24]: print('This is the best depth for Decision Tree Classifier: ', best_depth, '\nAcc
result = confusion_matrix(y_test.argmax(axis=1), y_predict_test.argmax(axis=1),)
print("Confusion Matrix:", "\n", result)
report = classification_report(y_test.argmax(axis=1), y_predict_test.argmax(axis=1),)
print("Classification Report:" , "\n", report)
```

This is the best depth for Decision Tree Classifier: 10

Accuracy score is: 0.13142857142857142

Confusion Matrix:

```
[[25  0  0  0  0  0  0]
 [25  2  0  0  1  0  0]
 [15  0  6  0  0  6  0]
 [26  0  0  1  0  0  0]
 [19  0  0  0  5  0  0]
 [15  0  2  0  0  4  0]
 [19  0  0  0  0  0  4]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.17	1.00	0.30	25
1	1.00	0.07	0.13	28
2	0.75	0.22	0.34	27
3	1.00	0.04	0.07	27
4	0.83	0.21	0.33	24
5	0.40	0.19	0.26	21
6	1.00	0.17	0.30	23
accuracy			0.27	175
macro avg	0.74	0.27	0.25	175
weighted avg	0.75	0.27	0.24	175

KNN

```
In [25]: k = range(1,102,3)
testing_accuracy = []
training_accuracy = []
score = 0

for i in k:
    knn = KNeighborsClassifier(n_neighbors = i)
    knn.fit(x_train, y_train)

    y_predict_train = knn.predict(x_train)
    training_accuracy.append(accuracy_score(y_train, y_predict_train))

    y_predict_test = knn.predict(x_test)
    acc_score = accuracy_score(y_test,y_predict_test)
    testing_accuracy.append(acc_score)

    print(i)


    if score < acc_score:
        score = acc_score
        best_k = i

sns.lineplot(k, training_accuracy)
sns.scatterplot(k, training_accuracy)
sns.lineplot(k, testing_accuracy)
sns.scatterplot(k, testing_accuracy)
plt.legend(['training accuracy', 'testing accuracy'])
```

1
4
7
10
13
16
19
22
25
28
31
34
37
40
43
46
49
52
55
58
61
64
67
70
73
76
79

82
85
88
91
94
97
100

Out[25]: <matplotlib.legend.Legend at 0x29663b0af40>




```
In [26]: print('This is the best depth for Decision Tree Classifier: ', best_depth, '\nAcc
result = confusion_matrix(y_test.argmax(axis=1), y_predict_test.argmax(axis=1),)
print("Confusion Matrix:", "\n", result)
report = classification_report(y_test.argmax(axis=1), y_predict_test.argmax(axis=
print("Classification Report:" , "\n", report)
```

This is the best depth for Decision Tree Classifier: 10

Accuracy score is: 0.44

Confusion Matrix:

```
[[25  0  0  0  0  0  0]
 [28  0  0  0  0  0  0]
 [27  0  0  0  0  0  0]
 [27  0  0  0  0  0  0]
 [24  0  0  0  0  0  0]
 [21  0  0  0  0  0  0]
 [23  0  0  0  0  0  0]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.14	1.00	0.25	25
1	0.00	0.00	0.00	28
2	0.00	0.00	0.00	27
3	0.00	0.00	0.00	27
4	0.00	0.00	0.00	24
5	0.00	0.00	0.00	21
6	0.00	0.00	0.00	23
accuracy			0.14	175
macro avg	0.02	0.14	0.04	175
weighted avg	0.02	0.14	0.04	175

```
In [27]: x_train.shape
```

```
Out[27]: (525, 37500)
```

MLP

In [28]: *# define the keras model*

```
model = Sequential()

model.add(Dense(units= 64, kernel_initializer = 'uniform', activation = 'relu', i
model.add(Dense(units= 128, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dense(units= 64, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dense(units= 32, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dense(units = 7, kernel_initializer = 'uniform', activation = 'softmax'
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 64)	240064
dense_1 (Dense)	(None, 128)	8320
dense_2 (Dense)	(None, 64)	8256
dense_3 (Dense)	(None, 32)	2080
dense_4 (Dense)	(None, 7)	231
=====		
Total params: 2,418,951		
Trainable params: 2,418,951		
Non-trainable params: 0		

```
In [29]: # compile the keras model
model.compile(optimizer = "Adam", loss = 'categorical_crossentropy', metrics = [

# fit the keras model on the dataset
history = model.fit(x_train, y_train, batch_size = 20, epochs = 25)
```

```
Epoch 1/25
27/27 [=====] - 1s 17ms/step - loss: 1.8014 - accurac
y: 0.2419
Epoch 2/25
27/27 [=====] - 0s 16ms/step - loss: 1.5498 - accurac
y: 0.2952
Epoch 3/25
27/27 [=====] - 0s 17ms/step - loss: 1.4296 - accurac
y: 0.3695
Epoch 4/25
27/27 [=====] - 0s 17ms/step - loss: 1.3027 - accurac
y: 0.4552
Epoch 5/25
27/27 [=====] - 0s 16ms/step - loss: 1.2219 - accurac
y: 0.4990
Epoch 6/25
27/27 [=====] - 0s 16ms/step - loss: 1.2073 - accurac
y: 0.5219
Epoch 7/25
27/27 [=====] - 0s 16ms/step - loss: 1.1443 - accurac
y: 0.4743
Epoch 8/25
27/27 [=====] - 0s 16ms/step - loss: 0.9984 - accurac
y: 0.5695
Epoch 9/25
27/27 [=====] - 0s 16ms/step - loss: 0.9170 - accurac
y: 0.6095
Epoch 10/25
27/27 [=====] - 0s 16ms/step - loss: 0.8727 - accurac
y: 0.6476
Epoch 11/25
27/27 [=====] - 0s 16ms/step - loss: 0.7510 - accurac
y: 0.6800
Epoch 12/25
27/27 [=====] - 0s 16ms/step - loss: 0.6952 - accurac
y: 0.7200
Epoch 13/25
27/27 [=====] - 0s 16ms/step - loss: 0.6431 - accurac
y: 0.7333
Epoch 14/25
27/27 [=====] - 0s 17ms/step - loss: 0.5594 - accurac
y: 0.7486
Epoch 15/25
27/27 [=====] - 0s 16ms/step - loss: 0.6699 - accurac
y: 0.7562
Epoch 16/25
27/27 [=====] - 0s 17ms/step - loss: 0.7145 - accurac
y: 0.7352
Epoch 17/25
27/27 [=====] - 0s 16ms/step - loss: 0.5015 - accurac
y: 0.8057
```

```

Epoch 18/25
27/27 [=====] - 0s 17ms/step - loss: 0.4875 - accuracy: 0.8095
Epoch 19/25
27/27 [=====] - 0s 17ms/step - loss: 0.3819 - accuracy: 0.8610
Epoch 20/25
27/27 [=====] - 0s 17ms/step - loss: 0.2616 - accuracy: 0.8876
Epoch 21/25
27/27 [=====] - 0s 17ms/step - loss: 0.3829 - accuracy: 0.8571
Epoch 22/25
27/27 [=====] - 0s 17ms/step - loss: 0.3570 - accuracy: 0.8533
Epoch 23/25
27/27 [=====] - 0s 17ms/step - loss: 0.3109 - accuracy: 0.8971
Epoch 24/25
27/27 [=====] - 0s 18ms/step - loss: 0.2291 - accuracy: 0.9162
Epoch 25/25
27/27 [=====] - 0s 17ms/step - loss: 0.2503 - accuracy: 0.9048

```

```

In [30]: accuracy = model.evaluate(x_test, y_test, verbose=1)[1]
print("Test: accuracy = ",accuracy*100,"%")

```

```

6/6 [=====] - 0s 9ms/step - loss: 2.9224 - accuracy: 0.4229
Test: accuracy = 42.28571355342865 %

```

CNN

```

In [31]: x_train = x_train.reshape(x_train.shape[0], 125,100,3)
x_test = x_test.reshape(x_test.shape[0], 125, 100, 3)
print(x_train.shape)
print(x_test.shape)

```

```

(525, 125, 100, 3)
(175, 125, 100, 3)

```

```
In [32]: input_shape = (125, 100, 3)
num_classes = 7

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', padding = 'Same', input_shape=input_shape))
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', padding = 'Same'))
model.add(MaxPooling2D(pool_size = (2, 2)))
model.add(Dropout(0.16))

model.add(Conv2D(64, (3, 3), activation='relu', padding = 'same'))
model.add(Conv2D(64, (3, 3), activation='relu', padding = 'Same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.20))

model.add(Flatten())
model.add(Dense(256, activation='relu'))
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.3))
model.add(Dense(num_classes, activation='softmax'))
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 125, 100, 32)	896
conv2d_1 (Conv2D)	(None, 125, 100, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 62, 50, 32)	0
dropout (Dropout)	(None, 62, 50, 32)	0
conv2d_2 (Conv2D)	(None, 62, 50, 64)	18496
conv2d_3 (Conv2D)	(None, 62, 50, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 31, 25, 64)	0
dropout_1 (Dropout)	(None, 31, 25, 64)	0
flatten (Flatten)	(None, 49600)	0
dense_5 (Dense)	(None, 256)	12697856
dense_6 (Dense)	(None, 128)	32896
dropout_2 (Dropout)	(None, 128)	0
dense_7 (Dense)	(None, 7)	903
=====		
Total params: 12,797,223		
Trainable params: 12,797,223		

Non-trainable params: 0

```
In [33]: model.compile(optimizer= "Adam",
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])

learning_rate_reduction = ReduceLROnPlateau(monitor='val_accuracy',
                                             patience=4,
                                             verbose=1,
                                             factor=0.5,
                                             min_lr=0.00001)
```

In [34]: `#x_train, x_validate, y_train, y_validate = train_test_split(x_train, y_train, te`

```
# Reshape image in 3 dimensions (height = 100, width = 125 , canal = 3)
x_train = x_train.reshape(x_train.shape[0], *(125, 100, 3))
x_test = x_test.reshape(x_test.shape[0], *(125,100, 3))
#x_validate = x_validate.reshape(x_validate.shape[0], *(100, 125, 3))
```

```
tf.config.run_functions_eagerly(True)
```

```
model.fit(x_train, y_train, batch_size = 100, epochs = 20)
```

Epoch 1/20

6/6 [=====] - 36s 6s/step - loss: 2.7245 - accuracy: 0.1638

Epoch 2/20

6/6 [=====] - 35s 6s/step - loss: 1.8154 - accuracy: 0.2229

Epoch 3/20

6/6 [=====] - 35s 6s/step - loss: 1.6871 - accuracy: 0.2933

Epoch 4/20

6/6 [=====] - 35s 6s/step - loss: 1.6260 - accuracy: 0.3010

Epoch 5/20

6/6 [=====] - 35s 6s/step - loss: 1.5808 - accuracy: 0.2990

Epoch 6/20

6/6 [=====] - 35s 6s/step - loss: 1.5184 - accuracy: 0.3714

Epoch 7/20

6/6 [=====] - 35s 6s/step - loss: 1.4675 - accuracy: 0.4457

Epoch 8/20

6/6 [=====] - 35s 6s/step - loss: 1.3887 - accuracy: 0.4514

Epoch 9/20

6/6 [=====] - 35s 6s/step - loss: 1.3177 - accuracy: 0.4800

Epoch 10/20

6/6 [=====] - 35s 6s/step - loss: 1.1809 - accuracy: 0.5295

Epoch 11/20

6/6 [=====] - 35s 6s/step - loss: 1.1108 - accuracy: 0.5314

Epoch 12/20

6/6 [=====] - 35s 6s/step - loss: 1.0377 - accuracy: 0.5657

Epoch 13/20

6/6 [=====] - 35s 6s/step - loss: 0.9896 - accuracy: 0.6057

Epoch 14/20

6/6 [=====] - 35s 6s/step - loss: 0.8993 - accuracy: 0.6495

Epoch 15/20

```
6/6 [=====] - 35s 6s/step - loss: 0.8305 - accuracy: 0.6819
Epoch 16/20
6/6 [=====] - 35s 6s/step - loss: 0.7601 - accuracy: 0.7010
Epoch 17/20
6/6 [=====] - 35s 6s/step - loss: 0.7278 - accuracy: 0.7162
Epoch 18/20
6/6 [=====] - 35s 6s/step - loss: 0.5842 - accuracy: 0.7714
Epoch 19/20
6/6 [=====] - 35s 6s/step - loss: 0.5808 - accuracy: 0.7733
Epoch 20/20
6/6 [=====] - 35s 6s/step - loss: 0.5257 - accuracy: 0.7752
```

Out[34]: <keras.callbacks.History at 0x2966748f880>

```
In [35]: test_loss, test_acc = model.evaluate(x_test, y_test)
print("Test Accuracy: ", test_acc*100, "%")
```

```
6/6 [=====] - 4s 579ms/step - loss: 1.4076 - accuracy: 0.5200
Test Accuracy: 51.99999809265137 %
```

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []: